

# Jinwoo Jeong

✉ [jwjeong@csl.korea.ac.kr](mailto:jwjeong@csl.korea.ac.kr)    [linkedin.com/in/jinwoojeong](https://www.linkedin.com/in/jinwoojeong)    [github.com/jinujeong](https://github.com/jinujeong)

## Research Interests

---

My research focuses on building efficient **serving systems** for deep neural networks and **large language models (LLMs)**, especially **memory management** and **serving scheduling** in GPU-centric environments.

◦ LLM Serving Systems   ◦ System for Machine Learning   ◦ Operating Systems   ◦ Computer Architecture

## Education

---

### Korea University

Seoul, South Korea

*Ph.D. Student in Electrical and Computer Engineering*

*Sep. 2024 – Feb. 2027*

◦ Advisor: Prof. Jeongseob Ahn

◦ Transferred from Ajou University, Aug. 2024

### Ajou University

Suwon, South Korea

*Ph.D. Student in Computer Engineering*

*Mar. 2020 – Aug. 2024*

### Ajou University

Suwon, South Korea

*B.E. in Software and Computer Engineering*

*Mar. 2016 – Feb. 2020*

## Publications

---

**Jinwoo Jeong** and Jeongseob Ahn. “An Efficient DNN Model Serving System using Layer-wise Caching and Direct-Host-Access”, *ACM Transactions on Computer Systems (TOCS)*, 2025.

**Jinwoo Jeong** and Jeongseob Ahn. “Accelerating LLM Serving for Multi-turn Dialogues with Efficient Resource Management”, *ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, 2025.

Woohyung Choi, **Jinwoo Jeong**, Hanhwi Jang, and Jeongseob Ahn. “GPU-centric Memory Tiering for LLM Serving with NVIDIA Grace Hopper Superchip”, *IEEE Computer Architecture Letters (CAL)*, 2025.

**Jinwoo Jeong**, Seungsu Baek, and Jeongseob Ahn. “Fast and Efficient Model Serving Using Multi-GPUs with Direct-Host-Access”, *ACM European Conference on Computer Systems (EuroSys)*, 2023.

Sangjin Choi, Taeksoo Kim, **Jinwoo Jeong**, Rachata Ausavarungrun, Myeongjae Jeon, Youngjin Kwon, and Jeongseob Ahn. “Memory Harvesting in Multi-GPU Systems with Hierarchical Unified Virtual Memory”, *USENIX Annual Technical Conference (ATC)*, 2022.

## Preprints

---

**Jinwoo Jeong**, Woohyung Choi, Myeongjae Jeon, and Jeongseob Ahn. “Taming Speculative Search for Test-Time Scaling in LLM Serving”, arXiv:2609.39334, 2026 (*Under Review*).

## Open-source Projects and Contributions

---

**SGLang** | <https://github.com/sgl-project/sglang>

◦ Contributed to an open-source LLM serving framework for efficient and scalable inference

**LMCache** | <https://github.com/LMCache/LMCache>

- Contributed to an open-source KV cache management system for LLM serving

**LMServe** | <https://github.com/Sys-KU/LMServe>

- Developed an open-source LLM serving framework that supports multi-level KV caching, request reordering, and disaggregated inference for efficient multi-turn LLM serving.

**DeepPlan** | <https://github.com/Sys-KU/DeepPlan>

- Released the EuroSys'23 artifact for multi-GPU DNN model serving, implementing layer-wise caching, direct-host-access, and parallel transmission to improve inference efficiency.

## Awards

---

|                                                                       |      |
|-----------------------------------------------------------------------|------|
| Best Paper Award Runner-up, <i>IEEE Computer Architecture Letters</i> | 2025 |
| Research Excellence Award, <i>Korea University</i>                    | 2025 |
| Gilles Muller Best Artifact Award, <i>EuroSys</i>                     | 2023 |

## Academic Services

---

Journal Reviewer

- 2026: IEEE Computer Architecture Letters

Shadow Program Committee

- 2026: EuroSys

## Technical Skills

---

**Languages:** C/C++, Python, Rust, CUDA, Triton, Go

**ML Frameworks:** SGLang, vLLM, LMCache, PyTorch, FlashInfer, FlashAttention

**Systems & Libraries:** NVIDIA UVM Driver, Linux Kernel, NIXL, NCCL, Asyncio, gRPC, Boost

**Developer Tools:** CMake/Make, Git, Docker, Nsight Systems/Compute, GDB, Perf, PyTest